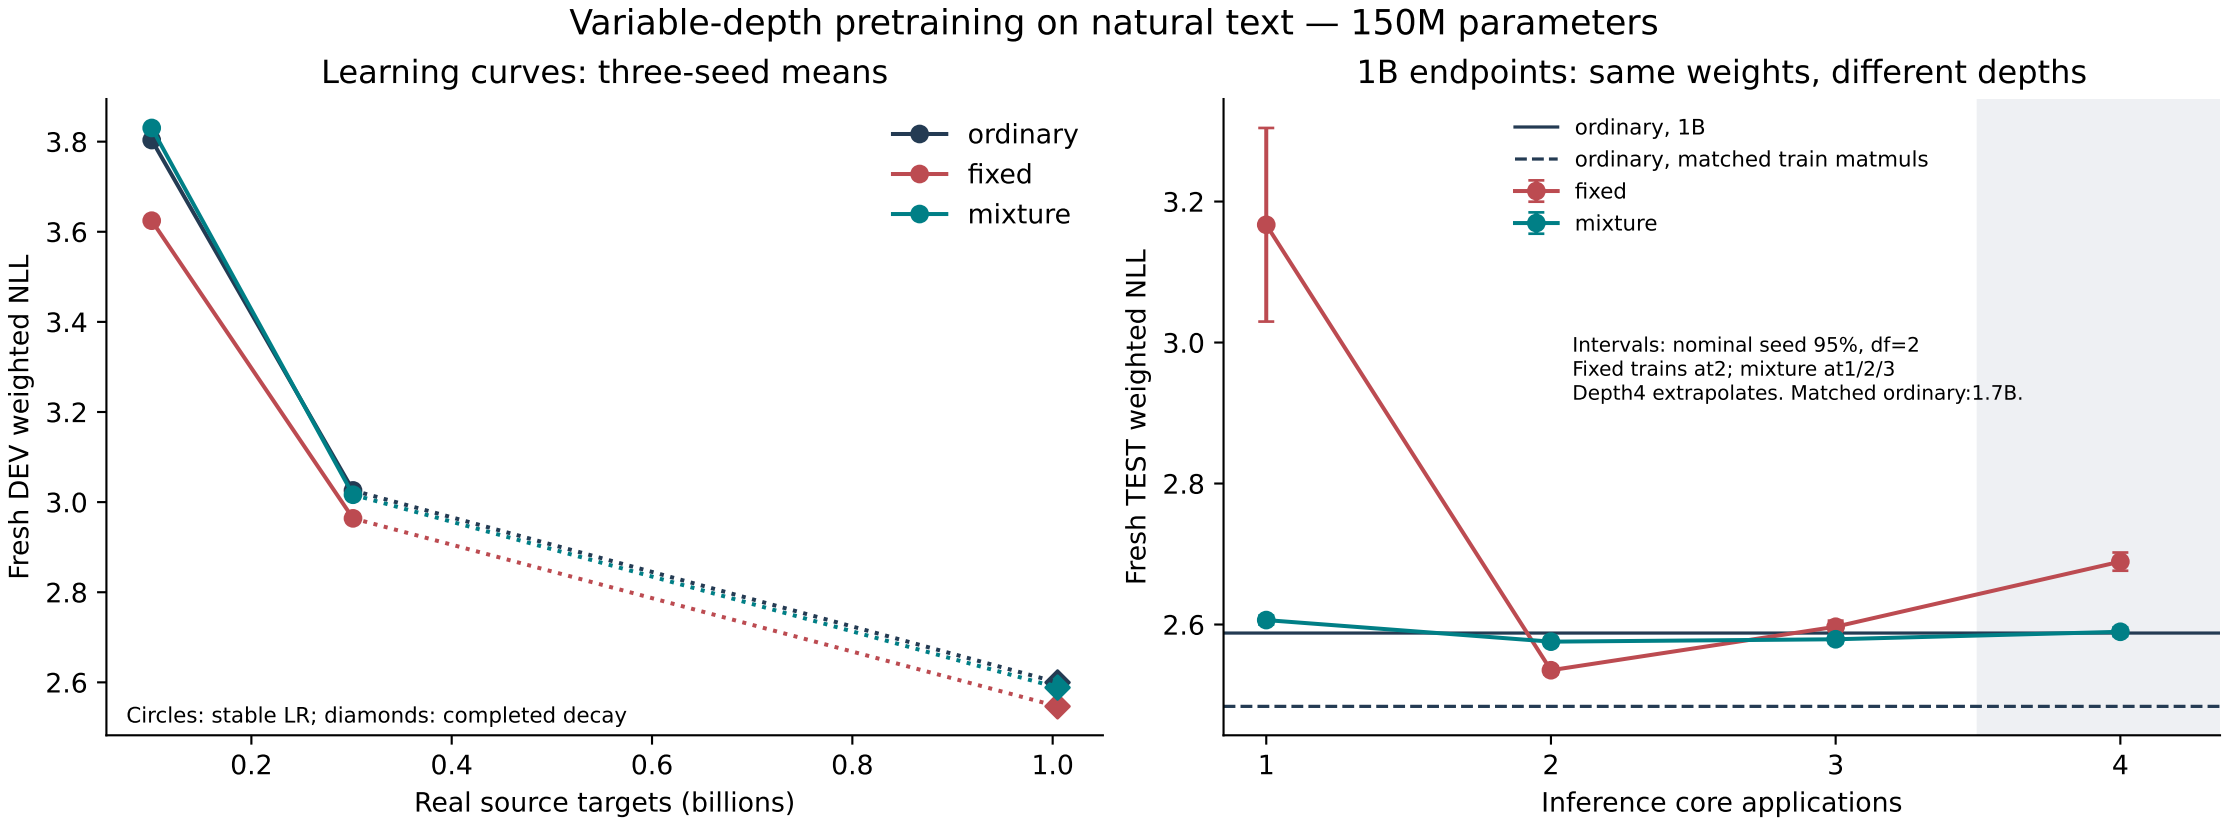
Variable-depth pretraining on natural text — 150M parameters
Learning curves: three-seed means
1B endpoints: same weights, different depths
ordinary
fixed
mixture
Fresh DEV weighted NLL
Real source targets (billions)
Circles: stable LR; diamonds: completed decay
ordinary, 1B
ordinary, matched train matmuls
fixed
mixture
Fresh TEST weighted NLL
Inference core applications
Intervals: nominal seed 95%, df=2
Fixed trains at2; mixture at1/2/3
Depth4 extrapolates. Matched ordinary:1.7B.